

# HPC Exam Project

Hybrid MPI/OpenMP 5-Point Stencil:  
Performance and Scalability Study

Giovanni Lucarelli

January 23, 2026



UNIVERSITÀ  
DEGLI STUDI  
DI TRIESTE

# Introduction

---

# Goals

---

1. **Optimize** the stencil method for the 2d heat equation
2. **Parallelize** using hybrid approach
3. Perform **scalability** study:
  - 3.1 Thread scaling
  - 3.2 Strong scaling
  - 3.3 Weak scaling

# Algorithm

## Heat equation (2d)

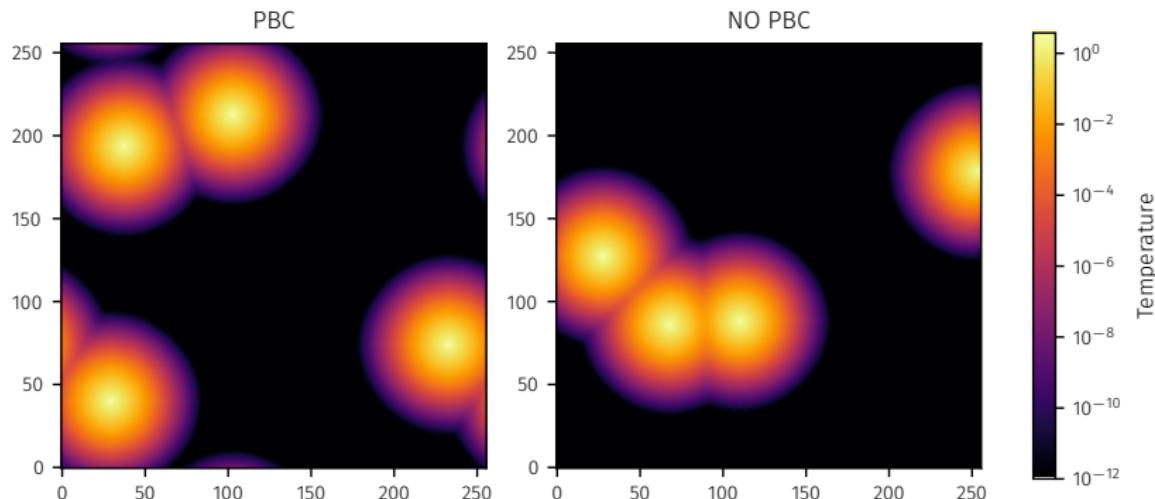
$$\partial_t u = \alpha(\partial_x^2 u + \partial_y^2 u)$$

## Finite difference integration

$$u_{i,j}^{(t+1)} = (1 - 4\alpha)u_{i,j}^{(t)} + \alpha \sum_{\langle i,j \rangle} u_{i,j}^{(t)}$$

$$x \in [0, L_x] \rightarrow i \in \{1, \dots, N_x - 1\}$$
$$y \in [0, L_y] \rightarrow j \in \{1, \dots, N_y - 1\}$$

# Code Correctness



# Build Configuration

---

```
1 CC      = mpicc
2 CFLAGS = -O3 -Wall -fopenmp -march=native
3 TARGET = stencil_mpi
4
5 ENABLE_OUTPUT ?= 0
```

# Parallelization: shared memory

## Implementation

```
1 #pragma omp parallel for schedule(static)
2 for (uint j = 1; j <= ysize; j++){
3     for ( uint i = 1; i <= xsize; i++){
4
5         // update rule
6
7     }
8 }
```

## Thread placement

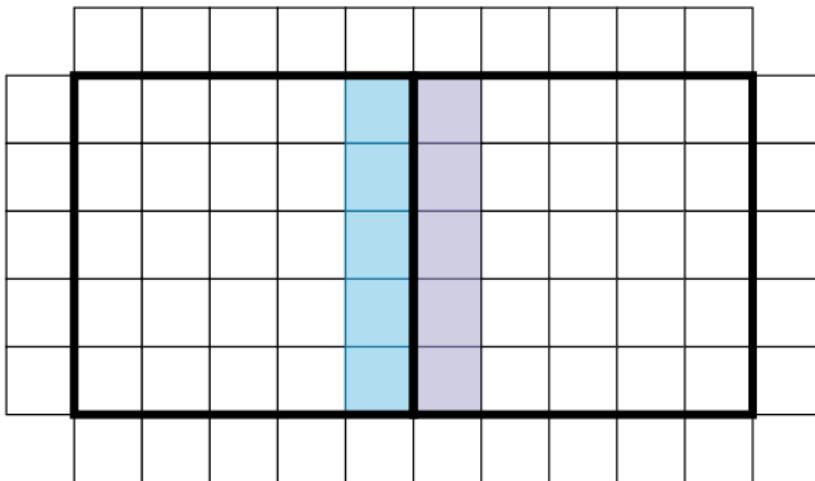
```
1 export OMP_PLACES=cores
2 export OMP_PROC_BIND=close
```

# First-touch

---



```
1 int memory_allocate ( ... ){
2
3     #pragma omp parallel for schedule(static)
4     for (int j = 0; j < Ny + 2; ++j){
5         for (int i = 0; i < Nx + 2; ++i) {
6             size_t idx = (size_t)j * (Nx + 2) + i;
7             planes_ptr[OLD].data[idx] = 0.0;
8             planes_ptr[NEW].data[idx] = 0.0;
9         }
10    }
11
12 }
```

# Parallelization: distributed memory



# Parallelization: distributed memory

right border **sends** to left halo  
left halo **receives** from right border



right halo **receives** from left border  
left border **sends** to right halo

# Parallelization: distributed memory

For each task:


```
1 // pack buffers
2
3 MPI_Irecv(...);
4
5 MPI_Isend(...);
6
7 update_internal();
8
9 MPI_Waitall();
10
11 // unpack buffers
12
13 update_border();
```

# Results

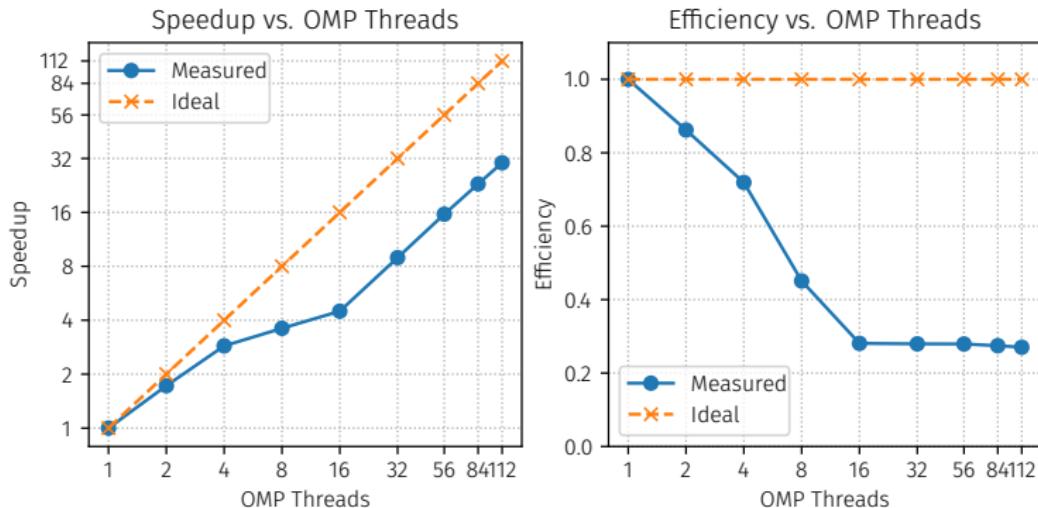
---

# Thread Scaling

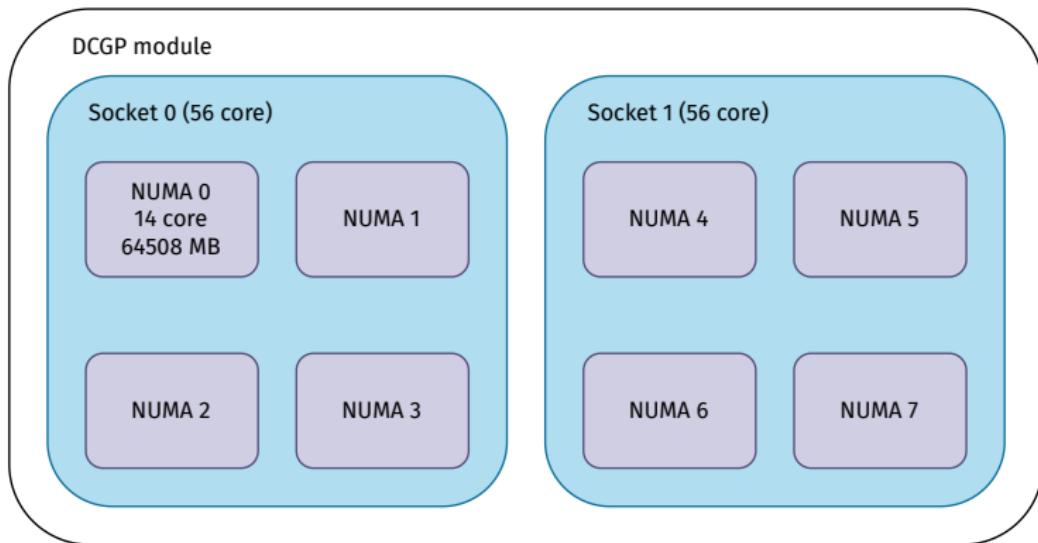
```
1 GRID_SIZE_X=16384  
2 GRID_SIZE_Y=16384  
3 N_STEPS=500  
4  
5 NODES=1  
6 N_TASKS_PER_NODE=1  
7 THREADS="1 2 4 8  
16 32 56 84  
112"
```



# Thread Scaling



# Node Architecture




# Node Architecture: distance matrix

```
1 [glucarel@lrdn4293 HPC-leonardo]$ numactl --hardware
2     available: 8 nodes (0-7)
3
4     node   0   1   2   3   4   5   6   7
5       0:  10  12  12  12  21  21  21  21
6       1:  12  10  12  12  21  21  21  21
7       2:  12  12  10  12  21  21  21  21
8       3:  12  12  12  10  21  21  21  21
9       4:  21  21  21  21  10  12  12  12
10      5:  21  21  21  21  12  10  12  12
11      6:  21  21  21  21  12  12  10  12
12      7:  21  21  21  21  12  12  12  10
```

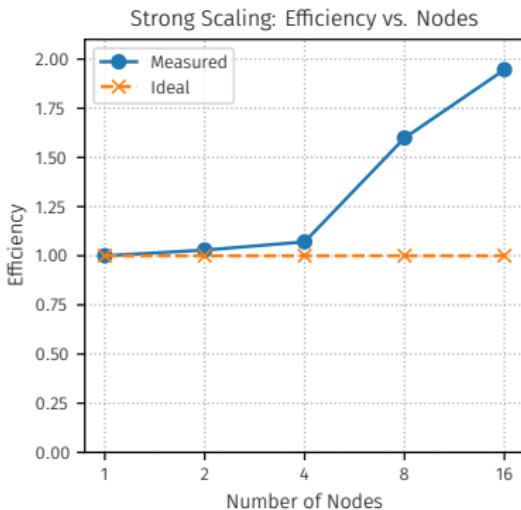
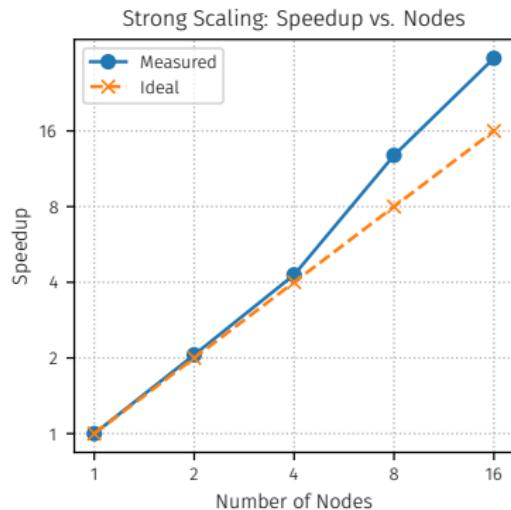
# Strong Scaling (1/4)

```
1 GRID_SIZE_X=16384  
2 GRID_SIZE_Y=16384  
3 N_STEPS=500  
4  
5 OMP_THREADS=14  
6 N_TASKS_PER_NODE=8  
7  
8 NODES="1 2 4 8 16"
```



Computation  $\gg$  Communication !

# Strong Scaling (1/4)



## Strong Scaling (1/4): Analysis

- 8 MPI tasks per node → one per **NUMA region**
- larger number of nodes → larger number of tasks  
(8 → 128) → smaller (local) grid size
- for 16 nodes (128 tasks as  $16 \times 8$ ), size  $N_x = N_y = 2^{14}$

$$mem_{128} = \frac{2^{14}}{16} \times \frac{2^{14}}{8} \times 16B \approx 33.5\text{MB}$$

```
1 [glucarel@lrdn4293 HPC-leonardo]$ lscpu | egrep 'L1d|L1i|L2|L3'
2     L1d cache:          48K
3     L1i cache:          32K
4     L2 cache:         2048K
5     L3 cache:        107520K
```

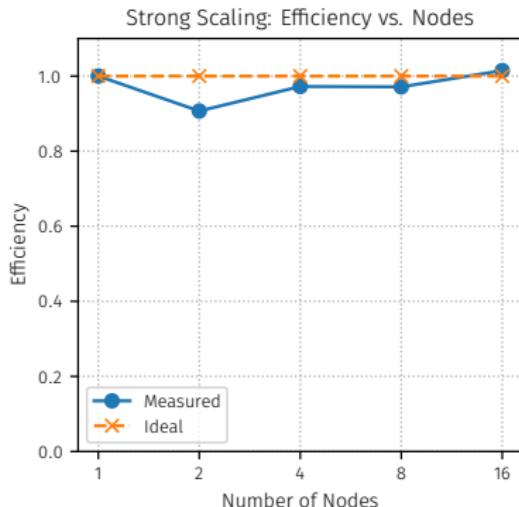
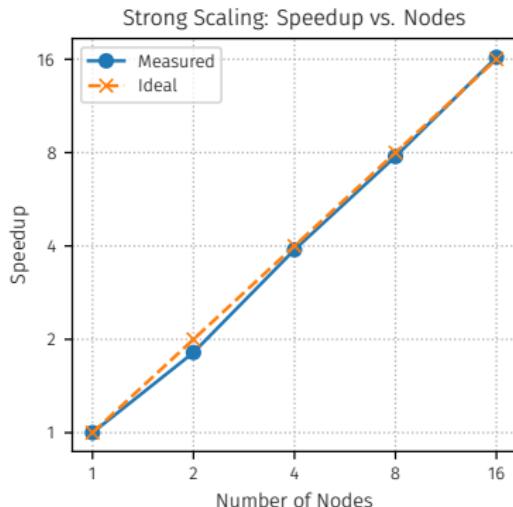
**Superlinearity** from cache effects exploitation!<sup>1</sup>

---

<sup>1</sup>If the grid is small enough

# Strong Scaling (2/4)

```
1 GRID_SIZE_X=16384  
2 OMP_THREADS=112  
3 N_TASKS_PER_NODE=1
```



## Strong Scaling (2/4): Analysis



---

- 1 MPI tasks per node → lower number of tasks  
(1 → 16)
- for 16 nodes (16 tasks as  $4 \times 4$ ), size  $N_x = X_y = 2^{14}$

$$mem_{16} = \frac{2^{14}}{4} \times \frac{2^{14}}{4} \times 16B \approx 268\text{MB}$$



# Strong Scaling (3/4)

```
1 GRID_SIZE_X=32768  
2 OMP_THREADS=14  
3 N_TASKS_PER_NODE=8
```



# Strong Scaling (4/4)

```
1 GRID_SIZE_X=65536  
2 OMP_THREADS=14  
3 N_TASKS_PER_NODE=8
```

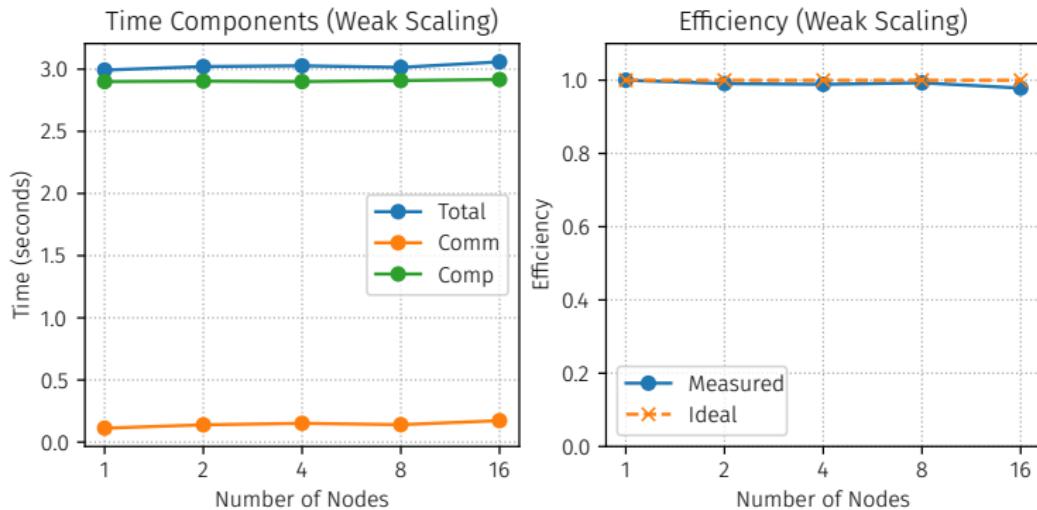


# Weak Scaling

---

```
1      LOCAL_X=4096
2      LOCAL_Y=4096
3      OMP_THREADS=14
4      TASKS_PER_NODE=8
5
6      for NODES in "1 2 4 8 16"; do
7          TOTAL_TASKS=$(( NODES * TASKS_PER_NODE ))
8
9          case "${TOTAL_TASKS}" in
10             8)    PX=4;   PY=2   ;;    # 1 node (8 ranks)
11             ...
12         esac
13
14         GRID_SIZE_X=$(( LOCAL_X * PX ))
15         GRID_SIZE_Y=$(( LOCAL_Y * PY ))
16         ...
17     done
```

# Weak Scaling



# Conclusion

---

## About stencil method:

- computation  $\gg$  communication
- data locality and effective cache usage improve the performance

## Possible improvements:

- MPI derived datatypes instead of buffers
- Higher-order stencil schemes

Thank You!