

HPC Exam Project

Scaling Study of the Stencil Method

Giovanni Lucarelli

October 20, 2025



UNIVERSITÀ
DEGLI STUDI
DITRIESTE

Introduction

Goals

1. **Optimize** the stencil method for the 2d heat equation
2. **Parallelize** using hybrid approach
3. Perform **scalability** study:
 - 3.1 Thread scaling
 - 3.2 Strong scaling
 - 3.3 Weak scaling

Algorithm

Heat equation (2d)

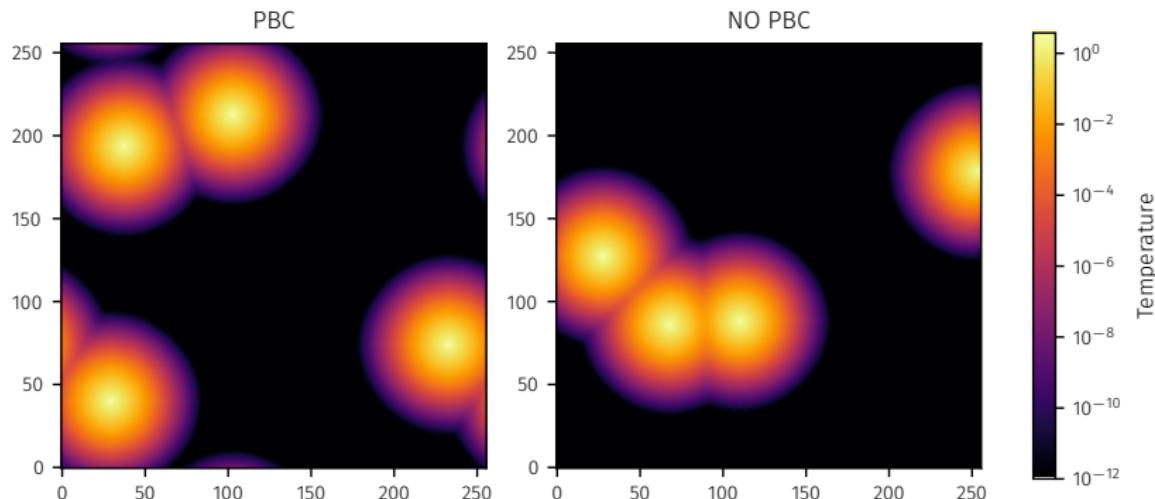
$$\partial_t u = \alpha(\partial_x^2 u + \partial_y^2 u)$$

Finite difference integration

$$u_{i,j}^{(t+1)} = (1 - 4\alpha)u_{i,j}^{(t)} + \alpha \sum_{\langle i,j \rangle} u_{i,j}^{(t)}$$

$$\begin{aligned} x \in [0, L_x] &\rightarrow i \in \{1, \dots, N_x - 1\} \\ y \in [0, L_y] &\rightarrow j \in \{1, \dots, N_y - 1\} \end{aligned}$$

Code Correctness



Optimization

- Compiler flags:

`-O3 -Wall -march=native`

- Preprocessor directive:

`#pragma GCC unroll`

Parallelization: shared memory

Implementation

```
1 #pragma omp parallel for schedule(static)
2 for (uint j = 1; j <= ysize; j++){
3     for ( uint i = 1; i <= xsize; i++){
4
5         // update rule
6
7     }
8 }
```

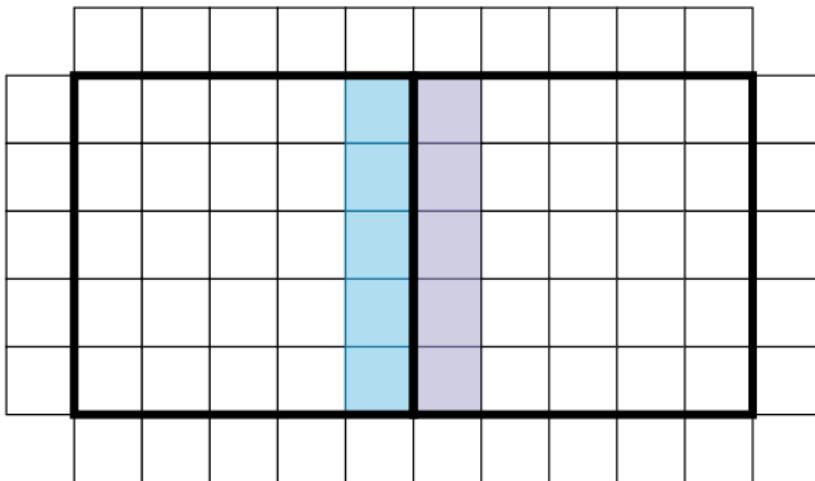
Thread placement and affinity

```
1 export OMP_PLACES=cores
2 export OMP_PROC_BIND=close
```

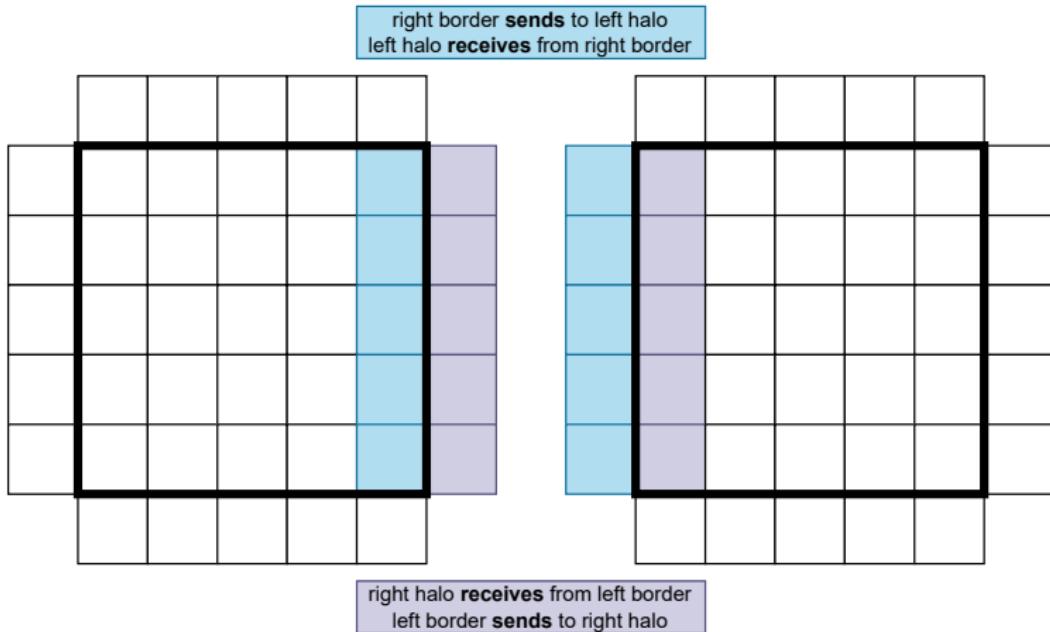
First-touch

```
1 int memory_allocate ( ... ){
2
3     #pragma omp parallel for collapse(2) schedule(static)
4     for (int j = 0; j < Ny + 2; ++j){
5         for (int i = 0; i < Nx + 2; ++i) {
6             size_t idx = (size_t)j * (Nx + 2) + i;
7             planes_ptr[OLD].data[idx] = 0.0;
8             planes_ptr[NEW].data[idx] = 0.0;
9         }
10    }
11
12 }
```

Parallelization: distributed memory



Parallelization: distributed memory



Parallelization: distributed memory

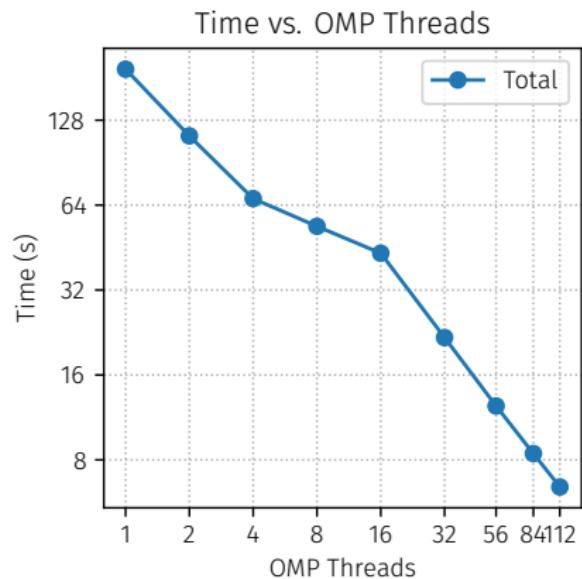
For each task:

```
1 // pack buffers
2
3 MPI_Irecv(...);
4
5 MPI_Isend(...);
6
7 update_internal();
8
9 MPI_Waitall();
10
11 // unpack buffers
12
13 update_border();
```

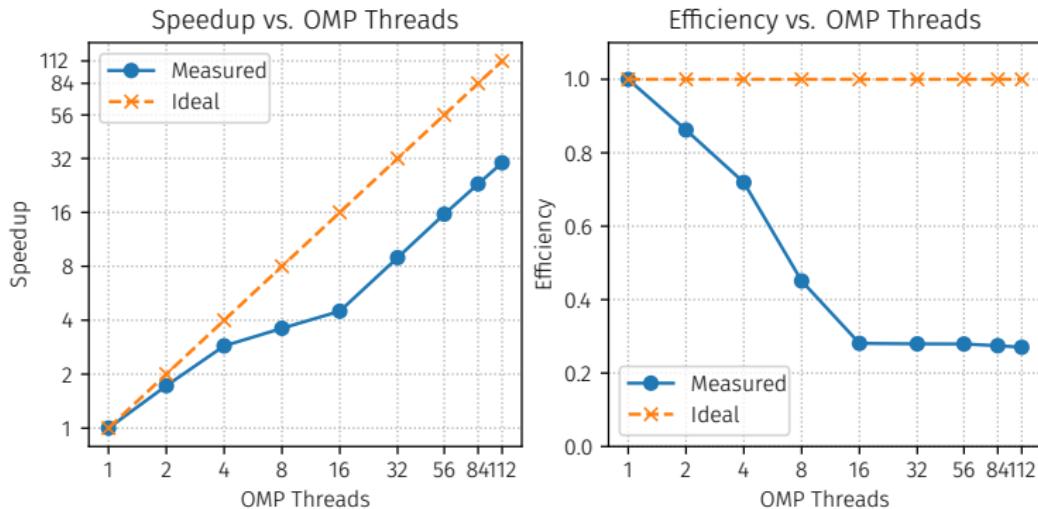
Results

Thread Scaling

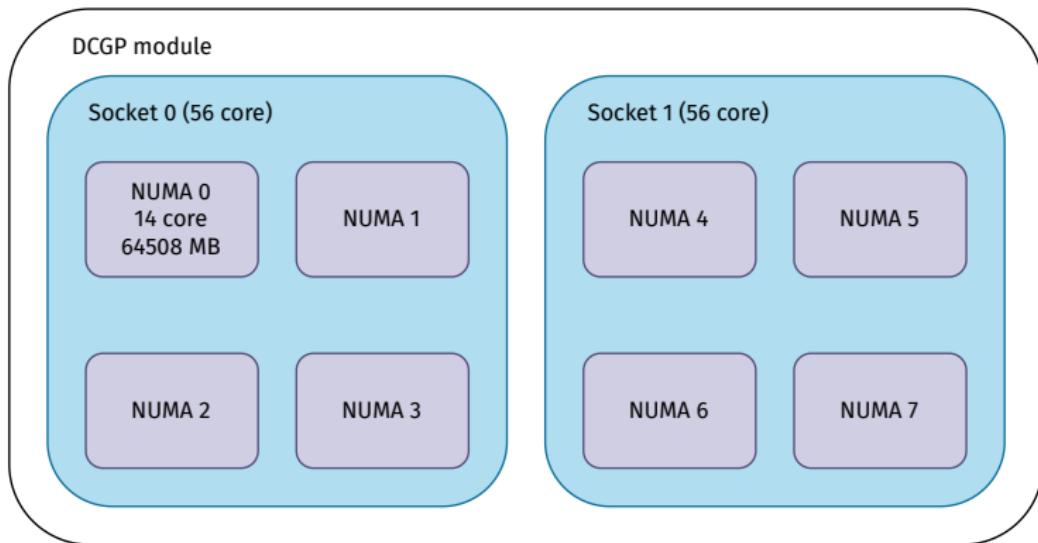
```
1   GRID_SIZE_X=16384  
2   GRID_SIZE_Y=16384  
3   N_STEPS=500  
4  
5   NODES=1  
6   N_TASKS_PER_NODE=1  
7   THREADS="1 2 4 8  
          16 32 56 84  
          112"
```



Thread Scaling



Node Architecture

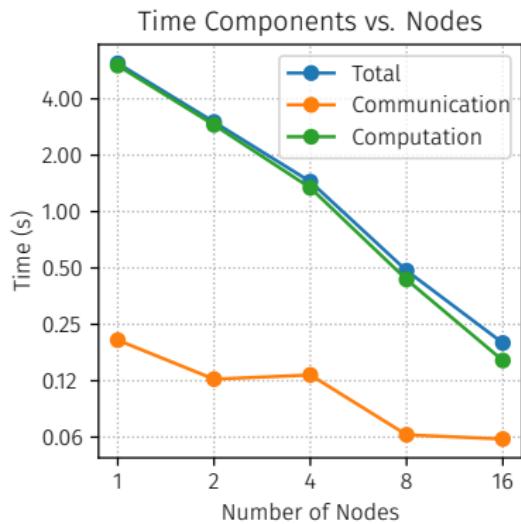


Node Architecture: distance matrix

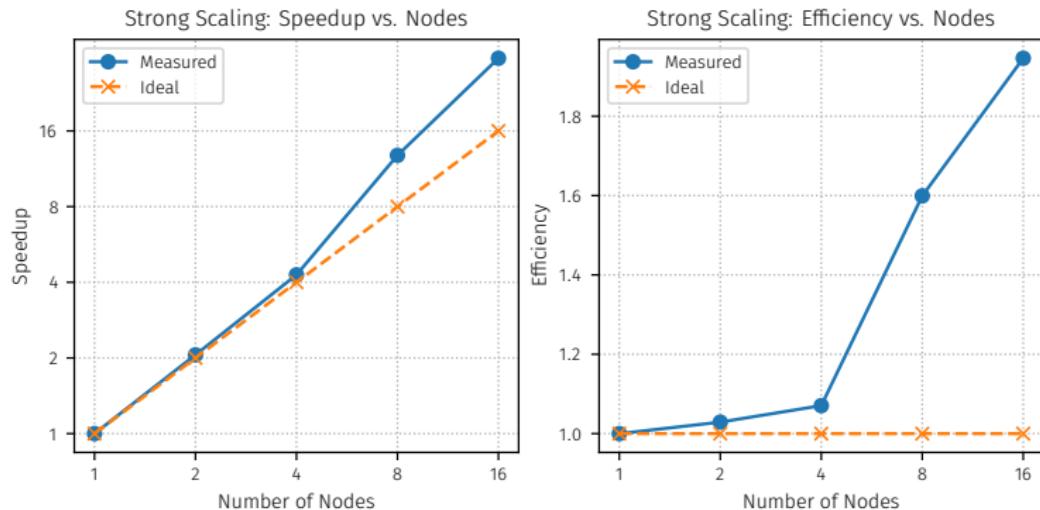
```
1 [glucarel@lrdn4293 HPC-leonardo]$ numactl --hardware
2     available: 8 nodes (0-7)
3
4     node   0    1    2    3    4    5    6    7
5       0:  10   12   12   12   21   21   21   21
6       1:  12   10   12   12   21   21   21   21
7       2:  12   12   10   12   21   21   21   21
8       3:  12   12   12   10   21   21   21   21
9       4:  21   21   21   21   10   12   12   12
10      5:  21   21   21   21   12   10   12   12
11      6:  21   21   21   21   12   12   10   12
12      7:  21   21   21   21   12   12   12   10
```

Strong Scaling (1/2)

```
1 GRID_SIZE_X=16384  
2 GRID_SIZE_Y=16384  
3 N_STEPS=500  
4  
5 OMP_THREADS=14  
6 N_TASKS_PER_NODE=8  
7  
8 NODES="1 2 4 8 16"
```



Strong Scaling (1/2)



Strong Scaling (1/2): Analysis

- 8 MPI tasks per node → one per **NUMA region**
- smaller grid as nodes increase (8 → 128)
- for nodes=16 (grid $N = 2^{14}$, 128 tasks = 16×8)


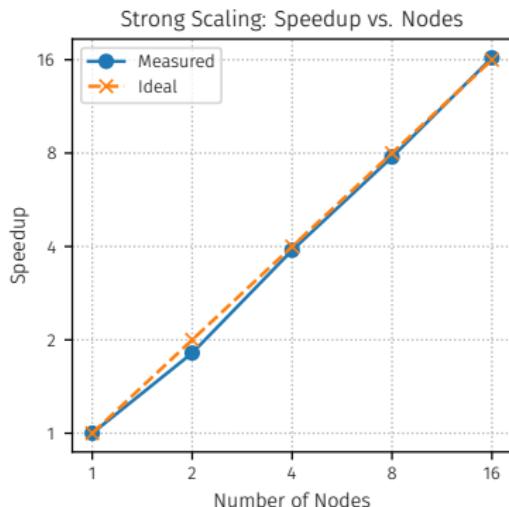
$$mem_{128} = 2^{10} \times 2^{11} \times 16B = 2^{25}B \approx 33.5\text{MB}$$

```
1 [glucarel@lrdn4293 HPC-leonardo]$ lscpu | egrep 'L1d|L1i|L2|L3'  
2     L1d cache:          48K  
3     L1i cache:          32K  
4     L2 cache:         2048K  
5     L3 cache:        107520K
```

Superlinearity comes from better cache effects
exploitation!

Strong Scaling (2/2)

1 OMP_THREADS=112
2 N_TASKS_PER_NODE=1



Strong Scaling (2/2): Analysis


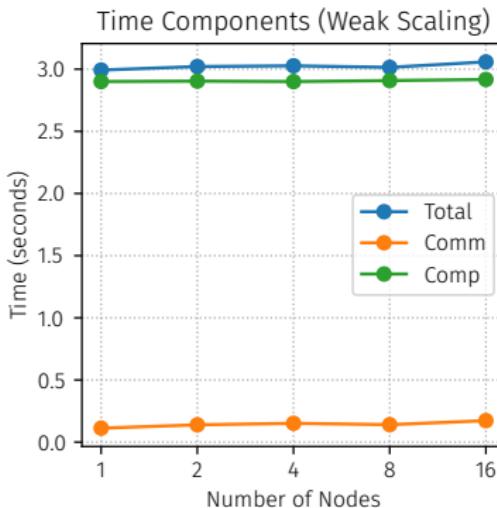
- 1 MPI tasks per node
- smaller grid as nodes increase ($1 \rightarrow 16$)
- for nodes=16 (grid $N = 2^{14}$, 16 tasks = 4×4)

$$mem_{16} = \frac{2^{14}}{4} \times \frac{2^{14}}{4} \times 16B = 2^{28}B \approx 268\text{MB}$$

Weak Scaling

```
1      LOCAL_X=4096
2      LOCAL_Y=4096
3      OMP_THREADS=14
4      TASKS_PER_NODE=8
5
6      for NODES in "1 2 4 8 16"; do
7          TOTAL_TASKS=$(( NODES * TASKS_PER_NODE ))
8
9          case "${TOTAL_TASKS}" in
10             8)    PX=4;   PY=2   ;;    # 1 node (8 ranks)
11             ...
12         esac
13
14         GRID_SIZE_X=$(( LOCAL_X * PX ))
15         GRID_SIZE_Y=$(( LOCAL_Y * PY ))
16         ...
17     done
```

Weak Scaling



Conclusion

About stencil method:

- computation » communication
- data locality and effective cache usage improve the performance

Possible improvements:

- Use MPI derived datatypes
-

Thank You!